

x86 Assembly Guide

[Contents: Registers | Memory and Addressing | Instructions | Calling Convention](#)

This guide describes the basics of 32-bit x86 assembly language programming, covering a small but useful subset of the available instructions and assembler directives. There are several different assembly languages for generating x86 machine code. The one we will use in CS216 is the Microsoft Macro Assembler (MASM) assembler. MASM uses the standard intel syntax for writing x86 assembly code.

The full x86 instruction set is large and complex (Intel's x86 instruction set manuals comprise over 2900 pages), and we do not cover it all in this guide. For example, there is a 16-bit subset of the x86 instruction set. Using the 16-bit programming model can be quite complex. It requires a segmented memory model, more restrictions on register usage, and so on. In this guide, we will limit our attention to more modern aspects of x86 programming, and delve into the instruction set only in enough detail to get a basic feel for x86 programming.

Resources

- [Guide to Using Assembly in Visual Studio](#) — a tutorial on building and debugging assembly code in Visual Studio
- [Intel x86 Instruction Set Reference](#)
- [Intel's Pentium Manuals](#) (the full gory details)

Registers

Modern (i.e. 386 and beyond) x86 processors have eight 32-bit general purpose registers, as depicted in Figure 1. The register names are mostly historical. For example, `eax` used to be called the accumulator since it was used by a number of arithmetic operations, and `ecx` was known as the counter since it was used to hold a loop index. Whereas most of the registers have lost their special purposes in the modern instruction set, by convention, two are reserved for segment pointers: the stack pointer (`esp`) and the base pointer (`ebp`).

For the `sx`, `tx`, `rx`, and `bx` registers, subregisters may be used. For example, the least significant 2 bytes of `ax` can be treated as a 16-bit register called `bx`. The least significant byte of `ax` can be used as a single 8-bit register called `bx`, while the most significant byte of `ax` can be used as a single 8-bit register called `bx`. These names refer to the same physical register. When a two-byte quantity is placed into `bx`, the update affects both `bx` and `bx`. These sub-registers are mainly hold-overs from older, 16-bit versions of the instruction set. However, they are sometimes convenient when dealing with data that are smaller than 32-bits (e.g. 1-byte ASCII characters).

When referring to registers in assembly language, the names are not case-sensitive. For example, the names `tx` and `tx` refer to the same register.

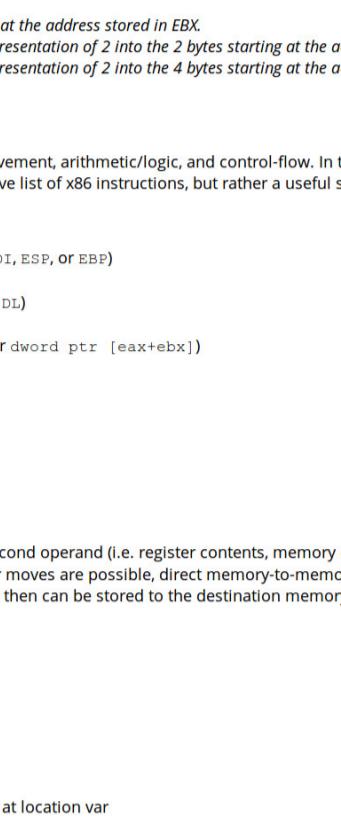


Figure 1. x86 Registers

Memory and Addressing Modes

Declaring Static Data Regions

You can declare static data regions (analogous to global variables) in x86 assembly using special assembler directives for this purpose. Data declarations should be preceded by the `.DATA` directive. Following this directive, the directives `DB`, `DW`, and `DD` are used to declare one, two, and four byte data locations, respectively. Declared locations can be labeled with names for later reference — this is similar to declaring variables by name, but abides by some lower level rules. For example, locations declared in sequence will be located in memory next to one another.

Example declarations:

```
.DATA
var1 DB 64          ; Declare a byte, referred to as location var, containing the value 64.
var2 DB ?           ; Declare an uninitialized byte, referred to as location var2.
DW ?               ; Declare a byte with no label, containing the value 10. Its location is var2 + 1.
DD 30000           ; Declare a 2-byte uninitialized value, referred to as location X.
Y DD 10 DUP(0)     ; Declare 10 4-byte words starting at location arr, all initialized to 0.
```

Unlike in high level languages where arrays can have many dimensions and are accessed by indices, arrays in x86 assembly language are simply a number of cells located contiguously in memory. An array can be declared by just listing the values, as in the first example below. Two other common methods used for declaring arrays of data are the `DUP` directive and the use of string literals. The `DUP` directive tells the assembler to duplicate an expression a given number of times. For example, `A DUP(2)` is equivalent to `2, 2, 2, 2`.

Some examples:

```
Z DD 1, 2, 3          ; Declare three 4-byte values, initialized to 1, 2, and 3. The value of location Z + 8 will be 3.
bytes DB 10 DUP(?)    ; Declare 10 uninitialized bytes starting at location bytes.
arr DD 100 DUP(0)      ; Declare 100 4-byte words starting at location arr, all initialized to 0.
str DB 'hello',0        ; Declare 6 bytes starting at the address str, initialized to the ASCII character values for hello and the null () byte.
```

Addressing Memory

Modern x86-compatible processors are capable of addressing up to 2³² bytes of memory. Memory addresses are 32-bits wide. In the examples above, where we used labels to refer to memory regions, these labels are actually replaced by the assembler with 32-bit quantities that specify addresses in memory. In addition to supporting referring to memory regions by labels (i.e. constant values), the x86 provides a flexible scheme for computing and referring to memory addresses: up to two of the 32-bit registers and a 32-bit constant can be added together to compute a memory address. One of the registers can be optionally pre-multiplied by 4 or 8.

The addressing modes can be used with many x86 instructions (we'll describe them in the next section). Here we illustrate some examples using the `mov` instruction that moves data between registers and memory. This instruction has two operands: the first is the destination and the second specifies the source.

Some examples of `mov` instructions using address computations are:

```
mov eax, [ebx]          ; Move the 4 bytes in the address contained in EBX into EAX
mov [eax], ebx          ; Move the contents of EBX into the 4 bytes of memory address var. Note, var is a 32-bit constant.
mov eax, [esi-4]         ; Move 4 bytes of memory address ESI-4 into EAX
mov [esi+eax], cl        ; Move the contents of CL into the byte at address ESI+EAX
mov edx, [esi+4*ebx]     ; Move the 4 bytes of data at address ESI+4*EBX into EDX
```

Some examples of invalid address calculations include:

```
mov eax, [ebx+ecx]       ; Can only add register values
mov [eax+esi+edi], ebx   ; At most 2 registers in address computation
```

Size Directives

In general, the intended size of the data item at a given memory address can be inferred from the assembly code instruction in which it is referenced. For example, in all of the above instructions, the size of the memory regions could be inferred from the size of the register operand. When we were loading a 32-bit register, the assembler could infer that the region of memory we were referring to was 4 bytes wide. When we were storing the value of a one byte register to memory, the assembler could infer that we wanted the address to refer to a single byte in memory.

However, in some cases the size of a referred-to memory region is ambiguous. Consider the instruction `mov [ebx], 2`. Should this instruction move the value 2 into the single byte at address `EBX`? Perhaps it should move the 32-bit integer representation of 2 into the 4 bytes starting at address `EBX`. Since either is a valid possible interpretation, the assembler must be explicitly directed as to which is correct. The size directives `BYTE PTR`, `WORD PTR`, and `DWORD PTR` serve this purpose, indicating sizes of 1, 2, and 4 bytes respectively.

For example:

```
mov BYTE PTR [ebx], 2      ; Move 2 into the single byte at the address stored in EBX.
mov WORD PTR [ebx], 2       ; Move the 16-bit integer representation of 2 into the 2 bytes starting at the address in EBX.
mov DWORD PTR [ebx], 2      ; Move the 32-bit integer representation of 2 into the 4 bytes starting at the address in EBX.
```

Instructions

Machine instructions generally fall into three categories: data movement, arithmetic/logic, and control-flow. In this section, we will look at important examples of x86 instructions from each category. This section should not be considered an exhaustive list of x86 instructions; but rather a useful subset. For a complete list, see Intel's instruction set reference.

We use the following notation:

<reg1>	Any 32-bit register (EAX, EBX, ECX, EDX, ESI, EDI, ESP, or EBP)
<reg1>	Any 16-bit register (AH, BH, CH, DH, AL, BL, CL, or DL)
<reg>	Any register
<mem>	A memory address (e.g., [eax], [var + 4], or dword ptr [eax+ebx])
<con32>	Any 32-bit constant
<con16>	Any 16-bit constant
<con8>	Any 8-bit constant
<con4>	Any 4-bit constant

Data Movement Instructions

mov — Move (OpCodes: 8B, 89, 8A, 8B, 8C, 8E, ...)

The `mov` instruction copies the data item referred to by its second operand (i.e. register contents, memory contents, or a constant value) into the location referred to by its first operand (i.e. a register or memory). While register-to-register moves are possible, direct memory-to-memory moves are not. In cases where memory transfers are desired, the source memory contents must first be loaded into a register, then can be stored to the destination memory address.

Syntax

```
mov <reg>,<reg>
mov <reg>,<mem>
mov <mem>,<reg>
mov <reg>,<const>
mov <mem>,<const>
```

Examples

```
mov eax, ebx — copy the value in ebx into eax
mov byte ptr [var], 5 — store the value 5 into the byte at location var
```

push — Push stack (OpCodes: FF, 89, 8A, 8B, 8C, 8E, ...)

The `push` instruction places its operand onto the top of the hardware-supported stack in memory. Specifically, `push` first decrements `ESP` by 4, then places its operand into the 32-bit location at address `[ESP]`. `ESP` (the stack pointer) is decremented by push since the `x86` stack grows down - i.e. the stack grows from high addresses to lower addresses.

Syntax

```
push <reg32>
push <mem>
```

Examples

```
push eax — push eax on the stack
push [var] — push the 4 bytes at address var onto the stack
```

pop — Pop stack

The `pop` instruction removes the 4-byte data element from the top of the hardware-supported stack in memory. It first moves the 4 bytes located at memory location `[ESP]` into the specified register or memory location, and then increments `ESP` by 4.

Syntax

```
pop <reg32>
pop <mem>
```

Examples

```
pop edi — pop the top element of the stack into EDI.
pop [var] — pop the top element of the stack into memory at location var.
```

lea — Load effective address

The `lea` instruction places the address specified by its second operand into the register indicated by its first operand. Note, the contents of the memory location are not loaded, only the effective address is computed and placed into the register. This is useful for obtaining a pointer into a memory region.

Syntax

```
lea <reg32>,<mem>
```

Examples

```
lea edi, [ebx+4*esi] — the quantity EBX+4*ESI is placed in EDI.
lea eax, [var] — the value in var is placed in EAX.
lea eax, [val] — the value val is placed in EAX.
```

Arithmetic and Logic Instructions

add — Integer Addition

The `add` instruction adds together its two operands, storing the result in its first operand. Note, whereas both operands may be registers, at most one operand may be a memory location.

Syntax

```
add <reg>,<reg>
add <reg>,<mem>
add <mem>,<reg>
add <reg>,<con>
add <mem>,<con>
```

Examples

```
add eax, 10 — EAX ← EAX + 10
add BYTE PTR [var], 10 — add 10 to the single byte stored at memory address var
```

sub — Integer Subtraction

The `sub` instruction stores the value of its first operand minus the result of subtracting the value of its second operand from the value of its first operand. As with `add`:

Syntax

```
sub <reg>,<reg>
sub <reg>,<mem>
sub <mem>,<reg>
sub <reg>,<con>
sub <mem>,<con>
```

Examples

```
sub eax, ebx — subtract ebx from eax. The result is stored in eax.
sub byte ptr [var], 5 — subtract 5 from the byte at location var.
```

inc, dec — Increment, Decrement

The `inc` instruction increments the contents of its operand by one. The `dec` instruction decrements the contents of its operand by one.

Syntax

```
inc <reg>
dec <reg>
```

Examples

```
inc ebx — increment ebx
dec ebx — decrement ebx
```

imul — Integer Multiplication

The `imul` instruction has two basic formats: two-operand (first two syntax listings above) and three-operand (last two syntax listings above).

Syntax

```
imul <reg32>,<reg32>
imul <reg32>,<mem>
imul <mem>,<reg32>
imul <reg32>,<const>
```

Examples

```
imul eax, ebx — multiply the contents of EAX by the 32-bit contents of the memory location var. Store the result in EAX.
```

idiv — Integer Division

The `idiv` instruction divides the contents of the 64-bit integer `EDX:EAX` (constructed by viewing `EDX` as the most significant four bytes and `EAX` as the least significant four bytes) by the specified operand value. The quotient result of the division is stored into `EAX`, while the remainder is placed in `EDX`.

Syntax

```
idiv <reg32>
idiv <mem>
```

Examples

```
idiv eax, ebx — divide the contents of EDX:EAX by the contents of EBX. Place the quotient in EAX and the remainder in EDX.
```

and, or, xor — Bitwise Logical and, or, and exclusive or

These instructions perform the specified logical operation (logical bitwise and, or, and exclusive or, respectively) on their operands, placing the result in the first operand location.

Syntax

```
and <reg>,<reg>
and <reg>,<mem>
and <mem>,<reg>
and <reg>,<con>
and <mem>,<con>
```

Examples

```
and eax, ebx — logical and of eax and ebx
and byte ptr [var], 0ffh — clear all but the last 4 bits of EAX.
```

xor

```
xor eax, ebx — set the contents of EDX to zero.
```

not

```
not <reg>
```

Examples

```
not eax — negate all bits in the byte at the memory location var.
```

neg

```
neg <reg>
```